

Risk Assessment of Aircraft for Commercial and Private Operations

1. Business Understanding

Project Goal

To analyze National Transportation Safety Board that includes aviation accident data from 1962 to 2023 about civil aviation accidents and selected incidents in the United States and international waters. To determine which aircraft are the lowest risk for the company, translate the findings into actionable insights that the head of the new aviation division can use to help decide which aircraft to purchase.

Data source

The NTSB aviation accident database contains information from 1962 and later about civil aviation accidents and selected incidents within the United States, its territories and possessions, and in international waters.

Key Questions:

1. Which aircraft type have the lowest accident rates?
2. Are there specific manufacturers associated with higher safety standards?
3. How do accident severities vary across aircraft types and manufacturers?
4. Are there specific times, places, or circumstances under which the risk is heightened for certain aircraft types?

Success Criteria:

1. Identify aircraft that have a low accident frequency and severity history
2. Provide actionable recommendations on model of aircraft that fit the company's operational needs.

2. Data Understanding

```
In [1]: ▶ #import libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```


```
In [2]: ▶ # Load the dataset.

df = pd.read_csv('AviationData.csv', engine='python')
df.head()
```

Out[2]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	Latitude	Longitude	Airport.Code
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States	NaN	NaN	NaN
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States	NaN	NaN	NaN
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States	36.922223	-81.878056	NaN
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States	NaN	NaN	NaN
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States	NaN	NaN	NaN

5 rows × 31 columns




In [3]: `# checking the tail of the data`
`df.tail()`

Out[3]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	Latitude	Longitude	Airport.Code	A
88884	20221227106491	Accident	ERA23LA093	2022-12-26	Annapolis, MD	United States	NaN	NaN	NaN	
88885	20221227106494	Accident	ERA23LA095	2022-12-26	Hampton, NH	United States	NaN	NaN	NaN	
88886	20221227106497	Accident	WPR23LA075	2022-12-26	Payson, AZ	United States	341525N	1112021W	PAN	
88887	20221227106498	Accident	WPR23LA076	2022-12-26	Morgan, UT	United States	NaN	NaN	NaN	
88888	20221230106513	Accident	ERA23LA097	2022-12-29	Athens, GA	United States	NaN	NaN	NaN	


5 rows × 31 columns




In [4]: `# Summary statistics of the numeric columns`
`df.describe()`

Out[4]:

	Number.of.Engines	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries	Total.Uninjured
count	82805.000000	77488.000000	76379.000000	76956.000000	82977.000000
mean	1.146585	0.647855	0.279881	0.357061	5.325440
std	0.446510	5.485960	1.544084	2.235625	27.913634
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000	0.000000	1.000000
75%	1.000000	0.000000	0.000000	0.000000	2.000000
max	8.000000	349.000000	161.000000	380.000000	699.000000

In [5]:  *# Size of the data*
df.shape

Out[5]: (88889, 31)

In [6]:  *# Summary of the Data*
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                    88889 non-null  object
2   Accident.Number                       88889 non-null  object
3   Event.Date                           88889 non-null  object
4   Location                             88837 non-null  object
5   Country                             88663 non-null  object
6   Latitude                             34382 non-null  object
7   Longitude                            34373 non-null  object
8   Airport.Code                         50249 non-null  object
9   Airport.Name                         52790 non-null  object
10  Injury.Severity                      87889 non-null  object
11  Aircraft.damage                      85695 non-null  object
12  Aircraft.Category                    32287 non-null  object
13  Registration.Number                  87572 non-null  object
14  Make                                88826 non-null  object
15  Model                               88797 non-null  object
16  Amateur.Built                       88787 non-null  object
17  Number.of.Engines                   82805 non-null  float64
18  Engine.Type                         81812 non-null  object
19  FAR.Description                     32023 non-null  object
20  Schedule                            12582 non-null  object
21  Purpose.of.flight                   82697 non-null  object
22  Air.carrier                         16648 non-null  object
23  Total.Fatal.Injuries                 77488 non-null  float64
24  Total.Serious.Injuries               76379 non-null  float64
25  Total.Minor.Injuries                 76956 non-null  float64
26  Total.Uninjured                     82977 non-null  float64
27  Weather.Condition                   84397 non-null  object
28  Broad.phase.of.flight                61724 non-null  object
29  Report.Status                       82508 non-null  object
30  Publication.Date                     75118 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.0+ MB
```

3. Data Preparation

Data Cleaning

```
In [7]: ▶ # Checking the number of missing values in each column  
df.isna().sum()
```

```
Out[7]: Event.Id                0  
Investigation.Type            0  
Accident.Number              0  
Event.Date                   0  
Location                     52  
Country                     226  
Latitude                    54507  
Longitude                   54516  
Airport.Code                 38640  
Airport.Name                 36099  
Injury.Severity              1000  
Aircraft.damage              3194  
Aircraft.Category            56602  
Registration.Number          1317  
Make                         63  
Model                        92  
Amateur.Built                102  
Number.of.Engines            6084  
Engine.Type                  7077  
FAR.Description              56000
```

```
In [8]: ▶ # I am going to drop columns that have roughly more than 25% of their data missing  
# More columns which are not important for my analysis, i drop for easier handling of the data  
columns_to_drop = ['Latitude', 'Longitude', 'Airport.Code', 'Accident.Number', 'Registration.Number',  
                   'Amateur.Built', 'Publication.Date', 'Report.Status', 'Engine.Type',  
                   'Airport.Name', 'Aircraft.Category', 'FAR.Description', 'Schedule',  
                   'Air.carrier', 'Broad.phase.of.flight']  
df_clean = df.drop(columns=columns_to_drop)
```

In [9]: `df_clean.columns`

Out[9]: Index(['Event.Id', 'Investigation.Type', 'Event.Date', 'Location', 'Country', 'Injury.Severity', 'Aircraft.damage', 'Make', 'Model', 'Number.of.Engines', 'Purpose.of.flight', 'Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured', 'Weather.Condition'], dtype='object')

In [10]: `# Rename some columns in the DataFrame`
`rename_columns = {`
 `'Event.Id': 'ID', 'Investigation.Type': 'Type', 'Event.Date': 'Date',`
 `'Injury.Severity': 'Injury_severity', 'Aircraft.damage': 'Damage_type',`
 `'Number.of.Engines': 'Engines', 'Purpose.of.flight': 'Flight_Purpose',`
 `'Total.Fatal.Injuries': 'Fatal_Injuries', 'Total.Serious.Injuries': 'Serious_Injuries',`
 `'Total.Minor.Injuries': 'Minor_Injuries', 'Total.Uninjured': 'Uninjured',`
 `'Weather.Condition': 'Weather'}`
`df_clean.rename(columns=rename_columns, inplace=True)`

In [11]: `df_clean.head()`

Out[11]:

	ID	Type	Date	Location	Country	Injury_severity	Damage_type	Make	Model	Engines	Flight_Purpose
0	20001218X45444	Accident	1948-10-24	MOOSE CREEK, ID	United States	Fatal(2)	Destroyed	Stinson	108-3	1.0	Perso
1	20001218X45447	Accident	1962-07-19	BRIDGEPORT, CA	United States	Fatal(4)	Destroyed	Piper	PA24-180	1.0	Perso
2	20061025X01555	Accident	1974-08-30	Saltville, VA	United States	Fatal(3)	Destroyed	Cessna	172M	1.0	Perso
3	20001218X45448	Accident	1977-06-19	EUREKA, CA	United States	Fatal(2)	Destroyed	Rockwell	112	1.0	Perso
4	20041105X01764	Accident	1979-08-02	Canton, OH	United States	Fatal(1)	Destroyed	Cessna	501	NaN	Perso

In [12]: `df_clean.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   ID                    88889 non-null  object 
 1   Type                  88889 non-null  object 
 2   Date                  88889 non-null  object 
 3   Location              88837 non-null  object 
 4   Country               88663 non-null  object 
 5   Injury_severity       87889 non-null  object 
 6   Damage_type           85695 non-null  object 
 7   Make                  88826 non-null  object 
 8   Model                 88797 non-null  object 
 9   Engines               82805 non-null  float64
10  Flight_Purpose          82697 non-null  object 
11  Fatal_Injuries        77488 non-null  float64
12  Serious_Injuries      76379 non-null  float64
13  Minor_Injuries        76956 non-null  float64
14  Uninjured             82977 non-null  float64
15  Weather               84397 non-null  object 
dtypes: float64(5), object(11)
memory usage: 10.9+ MB
```



```
In [13]: df_clean.isna().sum()
```

```
Out[13]: ID                0  
Type                0  
Date                0  
Location            52  
Country            226  
Injury_severity    1000  
Damage_type        3194  
Make                63  
Model              92  
Engines            6084  
Flight_Purpose        6192  
Fatal_Injuries     11401  
Serious_Injuries   12510  
Minor_Injuries     11933  
Uninjured          5912  
Weather            4492  
dtype: int64
```

```
In [14]: # Standardize the Injury_severity & Fatal_injuries
# The columns provide the same information
# Drop Fatal_Injuries column

df_clean.drop('Injury_severity', axis=1, inplace=True)
df_clean.head()
```

Out[14]:

	ID	Type	Date	Location	Country	Damage_type	Make	Model	Engines	Flight_Purpose	Fatal_Injuries
0	20001218X45444	Accident	1948-10-24	MOOSE CREEK, ID	United States	Destroyed	Stinson	108-3	1.0	Personal	2
1	20001218X45447	Accident	1962-07-19	BRIDGEPORT, CA	United States	Destroyed	Piper	PA24-180	1.0	Personal	4
2	20061025X01555	Accident	1974-08-30	Saltville, VA	United States	Destroyed	Cessna	172M	1.0	Personal	3
3	20001218X45448	Accident	1977-06-19	EUREKA, CA	United States	Destroyed	Rockwell	112	1.0	Personal	2
4	20041105X01764	Accident	1979-08-02	Canton, OH	United States	Destroyed	Cessna	501	NaN	Personal	1

Exploratory Data Analysis (EDA)

```
In [15]: # correlation
df_clean.corr()
```

Out[15]:

	Engines	Fatal_Injuries	Serious_Injuries	Minor_Injuries	Uninjured
Engines	1.000000	0.098505	0.046157	0.098162	0.406058
Fatal_Injuries	0.098505	1.000000	0.135724	0.073559	-0.015214
Serious_Injuries	0.046157	0.135724	1.000000	0.326849	0.052869
Minor_Injuries	0.098162	0.073559	0.326849	1.000000	0.147770
Uninjured	0.406058	-0.015214	0.052869	0.147770	1.000000

```
In [16]: # Plot a heatmap
plt.figure(figsize=(12,8))
sns.heatmap(df_clean.corr(),annot=True, cmap='coolwarm')
```

Out[16]: <AxesSubplot:>



Observations

1. Positive correlation: Engines and minor_injuries there is a strong positive correlation between these two variables by suggesting that as the number of engines increases the number of minor injuries tends to increase.
2. Negative correlations: Engines and Uninjured show a moderate negative correlation as the number of engines increases, number of uninjured individuals tends to decrease.


The heatmap suggests that there might be a relationship between the number of engine systems and likelihood of accidents and injuries.

```
In [23]: ▶ # Exploring the purpose of flights involved in accidents

top_10_purposes = df_clean['Flight_Purpose'].value_counts().nlargest(10).sort_values(ascending=False)

plt.figure(figsize=(6, 4))
top_10_purposes.plot(kind='bar')
plt.title('Distribution of Accidents by Flight Purpose (Top 10)')
plt.xlabel('Purpose of Flight')
plt.ylabel('Number of Accidents')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()

plt.show()
# Significant number of personal flights are responsible for aviation accidents
```



```
In [25]: # The relationship between Makes, Models and engine types with accidents
make_model_accident_counts = df_clean.groupby(['Make', 'Model']).size().reset_index(name='AccidentCount',
make_model_accident_counts = make_model_accident_counts.sort_values(by='AccidentCount', ascending=False)
make_model_accident_counts
```

Out[25]:

	Make	Model	AccidentCount
5745	Cessna	152	2168
5767	Cessna	172	1254
5811	Cessna	172N	996
15079	Piper	PA-28-140	812
5720	Cessna	150	716
...
8312	Engineering & Research	ERCOUPE 415-CD	1
8314	Engineering and Research	415C	1
8315	Engleman	PITTS S1	1
8316	English	PIETENPOL AIRCAMPER	1
20135	unknown	kit	1

20136 rows × 3 columns

```
In [28]: ▶ # Exploring what Make, Model and engine type involved in more accidents


df_clean['Engines'] = df_clean['Engines'].replace(0.0, np.nan) # some values in Engine column are 0. I am
# I am not sure if we have any aircraft with 0 engine. So I treat them as missing values

fig, axes = plt.subplots(1, 2, figsize=(12, 4))

top_10_make_model = make_model_accident_counts.head(10)
colors = sns.color_palette("dark", len(top_10_make_model))
axes[0].barh(top_10_make_model['Make'] + ' ' + top_10_make_model['Model'], top_10_make_model['AccidentCou
axes[0].set_xlabel('Number of Accidents')
axes[0].set_ylabel('Make and Model')
axes[0].set_title('Top 10 Aircraft Makes and Models Involved in Accidents')
axes[0].invert_yaxis()

axes[1].set_xticks([])
sns.countplot(data=df_clean, x='Engines', order=df_clean['Engines'].value_counts().index, ax=axes[1])
axes[1].set_title('Aircraft Engine Types Involved in Accidents')
axes[1].set_ylabel('Number of Accidents')
axes[1].set_xlabel('Engine Type')
axes[1].tick_params(axis='x', rotation=90)
axes[1].set_xlabel('Engine Type')

plt.show()
```



```
In [31]: ▶ # changing date type to the appropriate format and creating a column for seasons
df_clean['Date'] = pd.to_datetime(df_clean['Date'], format='%Y-%m-%d')
df_clean['Month'] = df_clean['Date'].dt.month
seasons = {
    12: 'Winter', 1: 'Winter', 2: 'Winter',
    3: 'Spring', 4: 'Spring', 5: 'Spring',
    6: 'Summer', 7: 'Summer', 8: 'Summer',
    9: 'Fall', 10: 'Fall', 11: 'Fall'
}

df_clean['Season'] = df_clean['Month'].map(seasons)
```

In [39]: *# Plot distribution of accidents by months and seasons*

```
fig, axes = plt.subplots(1, 2, figsize=(12, 4))

sns.countplot(data=df_clean, x='Month', palette='viridis', ax=axes[0])
axes[0].set_title('Distribution of Accidents by Month')
axes[0].set_xlabel('Month')
axes[0].set_ylabel('Number of Accidents')


month_order = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']
axes[0].set_xticks(range(12))
axes[0].set_xticklabels(month_order)

season_colors = {
    'Winter': 'blue',
    'Spring': 'red',
    'Summer': 'green',
    'Fall': 'yellow'
}

seasonal_accident_counts = df_clean['Season'].value_counts()

axes[1].bar(seasonal_accident_counts.index, seasonal_accident_counts, color=[season_colors.get(season, 't
axes[1].set_title('Distribution of Accidents by Season')
axes[1].set_xlabel('Season')
axes[1].set_ylabel('Number of Accidents')
plt.show()

# Most accidents happen in the summer
```

4. Recommendations

Personal flights are responsible for a high percentage of aviation accidents. In the area of personal flights, detailed education and training courses should be encouraged for pilots. The safety-first culture has to be nourished within the community of personal aviation. Minimizing the occurrence of accidents in personal flights depends on how safety is considered the priority in every personal flying activity. Whichever the weather condition, pilots must be very prepared and informed of the possible risks to be taken in personal aviation.

Most accidents occur in summer. In good weather, pilots may become overly confident and feel there is less risk than in poor weather conditions. This can result in a casual approach toward safety procedures and / or unsafe behavior, like flying low, going too fast, or using aerobatics, which raise the accident risk. The remedy is to encourage responsible flying practices and avoid taking unnecessary risks by proper training and awareness. It is also a re-emphasis on flying inside the safety envelope to avert accidents.

Aircraft makes and models, and engine types, in accidents. The Cessna 152 is the highest aircraft model in the number of accidents, followed by the Cessna 172. The opposite is observed in the case of the type of engine: the "1.0" type is dominant, while all others have much lower accident numbers. These findings might indicate that higher accident rates are related to models Cessna 152 and 172, combined with engine type "1.0.". Recommendations include: (1) Targeted safety investigations into these specific models and engine types to identify potential design flaws or operational issues; (2) Increased pilot training and education for these aircraft to address any recurring factors contributing to accidents; (3) Regular maintenance and inspections of these models and engines to ensure their continued airworthiness.